

SQLintersection

Eliminating Low Hanging Fruit

Jonathan Kehayias
Principal Consultant
SQLskills.com



SQL
intersection

Jonathan Kehayias



- Consultant/Trainer/Speaker
- Principal Consultant, [SQLskills.com](http://www.SQLskills.com)
 - Email: Jonathan@SQLskills.com
 - Blog: <http://www.SQLskills.com/blogs/jonathan>
 - Twitter: @SQLPoolBoy
- SQL Server MVP since October 2008
- Microsoft Certified Master: SQL Server 2008



Reminder: Intersect with Speakers and Attendees

- **Tweet tips and tricks that you learn and follow tweets posted by your peers!**
 - Follow: #SQLIntersection and/or #DEVIntersection
- **Join us – Wednesday Evening – for SQLafterDark**
 - Doors open at **7:00 pm**
 - Trivia game starts at **7:30 pm**
 - Winning team receives something fun!*
 - Raffle at the end of the night
 - Lots of great items to win including a seat in a SQLskills Immersion Event!*
 - The first round of drinks is sponsored by SentryOne and SQLskills



Overview

- **Row-By-Agonizing-Row processing**
- **Develop and test against realistic scale**
- **Profile your workload during development**
- **Use appropriate data types**
- **Sargability matters**

Row-By-AGONizing-Row Processing

- **Certain constructs force SQL Server into RBAR (row-by-agonizing-row) processing of results**
- **Well known:**
 - WHILE loops
 - Cursors
- **Less well known:**
 - Scalar user defined functions
 - Correlated subqueries

Cursors and Loops

- **Characteristics**

- Explicit cursor declaration
- WHILE loop
- SqlDataReader in application

- **Problems**

- Row based processing over Set Based

- **Replace with**

- Appropriate set based operation
- Move looping code into SQLCLR or Middle Tier
- Consume and dispose of SqlDataReader as quickly as possible

Correlated Sub-Queries

- **Characteristics**

- Refer to the outer query in the inner query in SELECT statement
- SELECT Statement used as column value in UPDATE

- **Problems**

- May cause row-by-row processing to occur
- Performance decreases exponentially as row count increases

- **Replace with**

- Table Join
- Derived Table Join
- Cross Joined Table Valued Function

Scalar User Defined Functions

- **Characteristics**

- Encapsulate common code blocks/business logic in a single call.
- If columns are passed as parameters it is not inline

- **Problems**

- Cause row-by-row processing to occur
- Performance decreases exponentially with data access

- **Replace with**

- Inline expressions
- Derived Table Join
- Cross Joined Inline Table Valued Function

Develop and Test Against Realistic Scale

- **Development databases often do not contain realistic datasets which can hide/mask potential performance problems**
 - Key Lookups on small data sets may become index scans on larger data sets
 - Missing index impacts may be hidden by data residing in memory for small data sets
 - RBAR problems are often hidden until data sizes scale up
- **Testing a single execution in isolation is not load testing**
 - Testing needs to be performed at scale through load generation to measure accumulated effects
 - Only testing at scale can identify “*death by 1000 cuts*” problems

Profile Your Workload During Development

- Learn to use Extended Events (2012+) or SQL Trace to profile your workload during testing
- Know the important events to watch for during development
 - Statement/Batch/RPC completed events
 - SP completed/Module End events – (procedure/trigger/function executions)
 - Execution warnings (sort, hash, missing join predicate)
- Profiling during development can uncover nasty RBAR issues and performance effecting side effects of trigger executions
- Be aware of “*observer overheads*” but not typically a problem with development/test workloads

Use Appropriate Data Types

- Understand the storage costs and implications of data types during schema design – especially for keys

| Data type | Range | Storage |
|-----------------|---|---------|
| bigint | -9,223,372,036,854,775,808 to 9,223,372,036,854,775,807 | 8 Bytes |
| int | -2,147,483,648 to 2,147,483,647 | 4 Bytes |
| smallint | -32,768 to 32,767 | 2 Bytes |
| tinyint | 0 to 255 | 1 Byte |

Use Appropriate Data Types (2)

- **Consider precision requirements for dates and times**
 - DATETIME2 = 6, 7, or 8 bytes with nanosecond precision
 - 6 bytes for precisions less than 3
 - 7 bytes for precisions 3 and 4
 - All other precisions require 8 bytes
 - DATETIME = 8 bytes with fractions of second precision
 - Fraction of seconds rounded to .000, .003, or .007 seconds
 - SMALLDATETIME = 4 bytes with minute precision
 - DATE = 3 bytes
- **Don't store dates or times as CHAR, VARCHAR, NCHAR, or NVARCHAR**

Sargability Matters

- A query is sargable (Search ARGument ABLE) if an index seek can be used to speed up the execution of the query
- Anti-patterns to sargable expressions include:
 - Functions in the WHERE clause
 - Implicit/Explicit data type conversions on a column
 - Leading wildcard expressions with LIKE '%<SearchTerm>'
 - Catch all queries and search procedures

Functions on WHERE Clause Columns

- **Characteristics**

- Used to change the data stored to match criteria being checked
- Conversion of data to a different type

- **Problems**

- Causes Table/Index Scan over Seek

- **Replace with**

- Appropriate Table Design to support business needs
- Indexed/Persisted Computed Column
- Indexed View
- Other coding paradigm to

Implicit/Explicit Column Conversions

- **Characteristics**

- Column data type is of lower precedence than filtering parameter / joining column data type
- Common in LINQ to SQL/EF and other ORMs

- **Problems**

- Causes Table/Index Scan over Seek

- **Replace with**

- Higher precedence column data type
- Matching data type for filtering parameter

Catch-All Search Queries

- **Characteristics**

- Used to search across multiple columns using parameters
- Not all parameters require input values
- WHERE clause similar to (`@Param1 IS NULL OR Column1 = @Param1`)

- **Problems**

- No optimized execution plan
- Causes Table/Index Scan over Seek

- **Replace with**

- Separate search procedures for different parameters passed
- Parameterized Dynamic SQL

Review

- **Row-By-Agonizing-Row processing**
- **Develop and test against realistic scale**
- **Profile your workload during development**
- **Use appropriate data types**
- **Sargability matters**

Questions?

SQLintersection

Don't forget to complete an online evaluation!

Eliminating Low Hanging Fruit

Your evaluation helps organizers build better conferences
and helps speakers improve their sessions.




SQL
intersection

Thank you!

Save the Date!

www.SQLintersection.com



2018

Mar 25-28

We're back in Orlando!



Leave the every day behind and enter a world of wonder and enchantment at the Walt Disney World® Resort. Located in the heart of the most magical place on earth, the Walt Disney World Swan and Dolphin Resort provides a truly extraordinary backdrop for our event! Beautiful tropical landscaping, tranquil waterways, and classic art and architecture work together to create a stunning landmark!